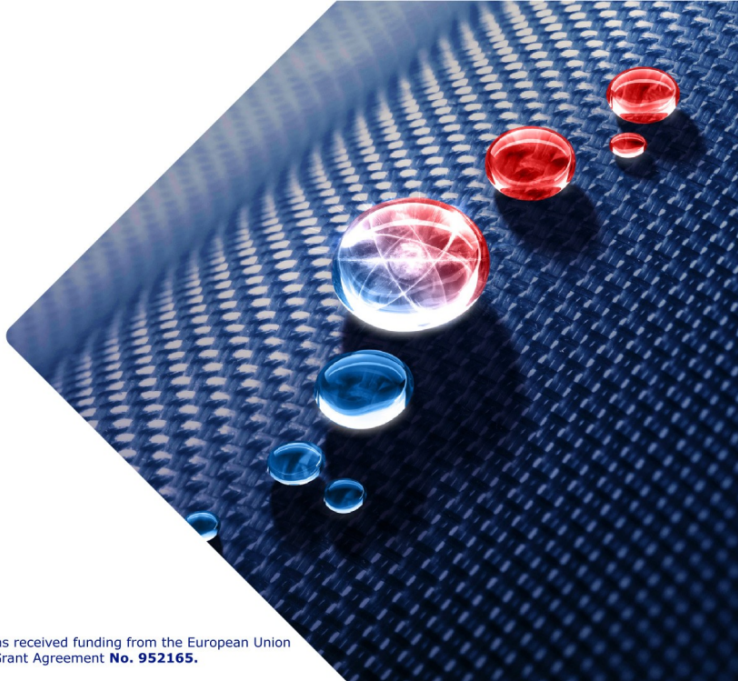


A Mixed-precision exploration of CHAMP

François Coppens

17/05/2021

Université de Versailles Saint-Quentin-en-Yvelines



Quick reminders

Language/compiler support

- C/C++/clang: LLVM versions 4.0 – 15.0
- Fortran/flang: LLVM versions ≤ 7.0

SSH connection to OLYMPE

```
> ssh -X <username>@olymp.e.calmip.univ-toulouse.fr
```

Minimal environment

- 1 > export MODULEPATH=\${MODULEPATH}:/usr/local/trex/modulefiles
- 2 > module load verificarlo

Verificarlo Singularity container

```
> sing-verificarlo
```

Verificarlo is a LLVM-based compiler extension. To this end Verificarlo offers compiler wrappers for the languages:

C `verificarlo-c`

C++ `verificarlo-c++`

Fortran `verificarlo-f`

To compile and link a single source file

```
verificarlo-<lang> source-file.<lang> -o <executable name>
```

Or compile and link several files

```
verificarlo-<lang> -c source-file1.<lang> -o source-file1.o
```

```
verificarlo-<lang> -c source-file2.<lang> -o source-file2.o
```

```
...
```

```
...
```

```
verificarlo-<lang> source-file1.o source-file2.o ... -o <executable>
```

To use the Verificarlo compiler and linker in Makefile or CMake projects we recommend setting the environment variables `CC`, `CXX` and `FC` for respectively the C, C++ and Fortran compiler.

These are usually respected by GNU Make and Kitware CMake, unless explicitly overridden by the user. To do this execute the following commands in the terminal

- `export CC=verificarlo-c`
- `export CXX=verificarlo-c++`
- `export FC=verificarlo-f`

Use OpenMPI wrapper scripts to include OpenMPI headers and libraries. Make sure to point these wrappers to the appropriate Verificarlo compiler wrappers by exporting the following environment variables

| | |
|---------------------|-----------------------------------|
| ■ export CC=mpicc | ■ export OMPI_C=verificarlo-c |
| ■ export CXX=mpic++ | ■ export OMPI_CXX=verificarlo-c++ |
| ■ export FC=mpifort | ■ export OMPI_FC=verificarlo-f |

This will ensure the compiler wrappers are chained in the correct order.



To use the Delta-Debug feature in Verificarlo the extra flag `--ddebug` needs to be passed to the compiler directly or as `CFLAGS/CXXFLAGS/FCFLAGS`

```
verificarlo-{c|c++|f} --ddebug \  
-c source-file.<lang> \  
-o source-file.o
```

or as

```
CFLAGS/CXXFLAGS/FCFLAGS += --ddebug
```




Mixed precision exploration in CHAMP using Delta-Debug

- Delta-Debug in Verificarlo
 - **Configurations** are the sets of floating-point instructions.
 - A **bug** is a numerical instability.
- Find unstable instructions for rounding / cancellations.
- Find instructions that can be run in lower precision.

| Step | Instructions with MCA noise | Numerically Stable |
|----------------|-----------------------------|--------------------|
| 1 | 1 2 3 4 | stable |
| 2 | 5 6 7 8 | unstable |
| 3 | 5 6 . . | stable |
| 4 | 7 8 | unstable |
| 5 | 7 . | unstable |
| Result (ddmin) | 7 . | |



Setting up the Delta-Debug helper scripts

- 1 Chose an observable: in this case it's the **total energy** and **one of the forces**
- 2 Prepare **ddRun**: script that runs the code and extracts the observable to monitor from the output
- 3 Prepare **ddCmp**: script that sets the **precision threshold** from the reference value and returns:
 - **true/pass**: if precision of current value does **NOT EXCEED** maximum deviation
 - **false/fail**: if precision of current value **EXCEEDS** the maximum deviation
- 4 Prepare a small **Makefile**: to automate setting Verificarlo **backend**, **precision**, **# of ddebug runs** and **vfc_ddebug** invocation

```
OUTDIR=$1

BIN=./vfcdd_champ
INP=input.inp

for i in $(seq 1 1)
do
    $BIN -i $INP -o /tmp/res.dat 2> ${OUTDIR}/ddrun.error
    awk '/total E/{print$4}' /tmp/res.dat > $OUTDIR/res$i.dat
    rm /tmp/res.dat
done
```

```
import sys
import glob
import os
import numpy as np
import math

PRECISION_THRESHOLD = 5 # number of DECIMAL digits
REFDIR = sys.argv[1]
CURRDIR = sys.argv[2]
res = np.zeros(shape=0)
```

```
def read_output(DIR, RES):  
    resList = glob.glob(os.path.join(DIR, 'res*.dat'))  
    resList += glob.glob(os.path.join(REFDIR, 'res*.dat'))  
    for resFile in resList:  
        with open(resFile) as f:  
            RES = np.append(RES, float(f.read()))  
    return RES
```

```
res = read_output(CURRDIR, Res)

s = math.log2(abs((res[0]-res[1])/res[1]))/math.log2(10)
    # see next slide

with open("{}res.stat".format(CURRDIR), 'w') as f:
    print(f"Stat. file loc = {CURRDIR}/res.stat")
    f.write("s = {}\n".format(s))

sys.exit(1 if s < PRECISION_THRESHOLD else 0)
```


Precision in number of bits

$$s_2 = -\log_2 \left| \frac{x_{\text{VPREC}} - x_{\text{IEEE}}}{x_{\text{IEEE}}} \right|$$

Precision in number decimals

$$s_{10} = \frac{s_2}{\log_2(10)}$$

```
NRUNS=1
PRECISION=24
BACKEND="libinterflop_vprec.so --precision-binary64=${PRECISION}"

dd: vfcdd_champ
    rm -rvf dd.line/
    INTERFLOP_DD_NRUNS=${NRUNS} VFC_BACKENDS=${BACKEND} \
    vfc_ddebug ddRun_vp ddCmp_vp

dderrors: dd.line/rddmin-cmp/dd.line.exclude
    bash -c "vim -q <(<./vfc_dderrors.py ./vmc $<)"
```



After running make dd you should see output like

```
INTERFLOP_DD_NRUNS=1 VFC_BACKENDS="libinterflop_vprec.so \  
  --precision-binary64=17" vfc_ddebug ddRun_vp ddCmp_vp  
dd.line/8a391a01211abd91e3d54cfeca2897f4 --( run )-> FAIL(0)  
dd.line/8a391a01211abd91e3d54cfeca2897f4 --(cache) -> FAIL  
dd.line/525c227b98a0ae2e100829a4100dbcf4 --( run )-> FAIL(0)  
dd.line/4ebd0a8a27fa25253644536b5ea1cbe8 --( run )-> FAIL(0)  
dd.line/8a8069fe0ee2b665eab4c7034fba118d --( run )-> PASS(+1->1)  
dd.line/b434b51d52907341f6f05a9e4128d8e1 --( run )-> FAIL(0)  
dd.line/53769469e1628356f7b2aa219c9df18a --( run )-> FAIL(0)  
dd.line/d2d274f8db7675ba550c01ce6979c32a --( run )-> PASS(+1->1)  
dd.line/93e0258c4d272730e115c64951366207 --( run )-> FAIL(0)  
dd.line/c7d60fc315f6b8e0ec55cb39c6a4edd1 --( run )-> PASS(+1->1)  
dd.line/91384f0e34618621669fdd58e04583c3 --( run )-> FAIL(0)  
ddmin0 (0x0000000000521ced: splfit_ at splfit.f:47):
```

```
drwxr-xr-x. coppens p22064 04bb7a6780419ee429cd806fe66cce66
drwxr-xr-x. coppens p22064 05928b62472d3723566e3b4d193bcb2c
drwxr-xr-x. coppens p22064 0598f288b121391b3cd8f9f2e948c8d2
drwxr-xr-x. coppens p22064 d914463b3d0055e77e0128b74ceb5291
lrwxrwxrwx. coppens p22064 ddmin0 -> 05928b62472d3723566e3b4d193bcb2c
lrwxrwxrwx. coppens p22064 ddmin1 -> d914463b3d0055e77e0128b74ceb5291
drwxr-xr-x. coppens p22064 dee78da1159ba7d45cfd5809f9799e22
drwxr-xr-x. coppens p22064 df02513b130ff37c200ebf49f7aa22a7
drwxr-xr-x. coppens p22064 fd5e7760f41617300b4e5b49b6a55843
drwxr-xr-x. coppens p22064 ff680b1e38e2818b5e9f67029f123644
lrwxrwxrwx. coppens p22064 rddmin-cmp -> df02513b130ff37c200ebf49f7aa22a7
drwxr-xr-x. coppens p22064 ref
```

```
ff680b1e38e2818b5e9f67029f123644
```

```
ref
```

```
|-- dd.line.exclude
```

```
|-- checkRef.err
```

```
|-- dd.line.include
```

```
|-- checkRef.out
```

```
`-- dd.run1
```

```
|-- dd.err
```

```
    |-- dd.compare.err
```

```
|-- dd.line
```

```
    |-- dd.compare.out
```

```
|-- dd.line.%%p
```

```
    |-- dd.run.err
```

```
|-- dd.out
```

```
    |-- ddrun.error
```

```
|-- ddrun.error
```

```
    |-- dd.run.out
```

```
|-- res1.dat
```

```
    |-- res1.dat
```

```
`-- res.stat
```

```
    |-- res.stat
```

```
`-- returnVal
```

When done, minimal set can be found in `$PWD/dd.line/rddmin-cmp`. It contains the files

- `dd.line.exclude`: contains the functions/subroutines that break at the chosen precision.
- `dd.line.include`: contains the functions/subroutines that can be changed to the lower precision without affecting the precision of the total energy

Table: CHAMP/VMC Butadiene CIPSI. Investigated precision on TOTAL ENERGY for all functions.

| Function name | Required precision (VPREC) |
|---------------|----------------------------|
| splfit | Double |
| ALL OTHERS | Single |

Conclusion: When only interested in the energy we can run most of CHAMP at single precision, potentially gaining a significant speedup.

Table 2: CHAMP/VMC Butadiene CIPSI. Investigated precision on FORCES for functions of runtime $\geq 5\%$ of total run-time.

| Function name | Time spent (%) | | Required precision (VPREC) |
|------------------------|----------------|----------|----------------------------|
| | 500 dets | 15k dets | |
| orbitals | 22.02 | 5.71 | Double |
| nonloc | 11.7 | 3.15 | |
| > orbitals_quad:395 | 7 | | Single |
| orbitalse | 5.56 | 3.36 | Not called in VFC-DD |
| optjas.deloc | 4.66 | 16.45 | Singel |
| splfit | 4 | | Double |
| multideterminante_grad | 3.62 | 2.48 | |
| multideterminante | 3.59 | 14.86 | Double |
| multideterminant_hpsi | 3.56 | 3.49 | |
| n0_inc | 2.6 | | |
| basis_fns_vgl | 2.48 | | |
| basis_fnse_v | 2.43 | | |
| optorb.compute | 2.08 | | |
| compute_ymat | | 12.8 | Double |
| detsav | | 4.66 | Not called in VFC-DD |
| __powr8i4 | 7.16 | 1.88 | |
| __libm_log_l9 | 2.43 | 0.58 | |

Conclusion: sometimes life is more complicated

MPI flang-7 is not able to assign multiple file descriptors to the same file →
prevents from running more than one MPI process!

MORE PROBLEMS TO BE ADDED